

Recommendations

Final Report

Holubiec, Maciej
Romero, Maria Jimena
Von Chamier, Paul

1. Problem Statement and Motivation

The project's objective is to create a recommendation system for restaurants using regressions. We will use regressions to predict ratings (number of stars) per user and business. In specific, we will make two Baselines, one with sample averages and another with a Regularized Regression, which we will compare to our Matrix Factorization, according to the Root Mean Square Error. The data used comes from Yelp reviews, where we can identify users and business, among other characteristics like type of business, geographic region, among others.

2. Introduction and Description of Data

The project is relevant in the consumption industry, as currently modern consumers are inundated with choices, and matching consumers with the most appropriate products is key to enhancing user satisfaction and loyalty¹. However, there are more sectors that can benefited from this method. In public policy, for instance, there is wide potential to better match citizens preferences to services and policies. Citizens can benefit from better information on school, public health locations, doctors, among other. Also, the Matrix Factorization can be used not only when dealing with consumers, but also when trying to fill any missing values between two different parties (like even implied costs between companies).

The challenge comes from the fact that on average, users rate only a few places and places are rated by only a few users. According to the EDA, we found out that we see that the average rating issued per user 2.5 and that the average number of ratings received per business is 12.5. This large amount of missing values can create difficulties when doing regressions, and also create a matrix with more number of variables than observations. To fix this, we might only select the most frequented restaurants and users.

The Yelp data is distributed among 5 different json files, out of which 3 are being used: business.json, review.json and user.json. In total, we started with 4.7 MM reviews. To build our models, we will need to transform review database to create a user_id - business_id matrix (filled with ratings). With our baseline model, we will fill the missing data with averages or predictions from a multiple regression. With the latent matrix model, we will try to predict using matrix factorization backed up by alternating least squares regularization.

3. Literature Review/Related Work

Matrix factorization for recommender system is a recent method that became popular in 2006 when Netflix launched a competition to improve its recommendations.² During the following years, Matrix Factorization methods were the most important class of techniques for winning the Netflix Prize³.

¹ Matrix factorization for recommender systems

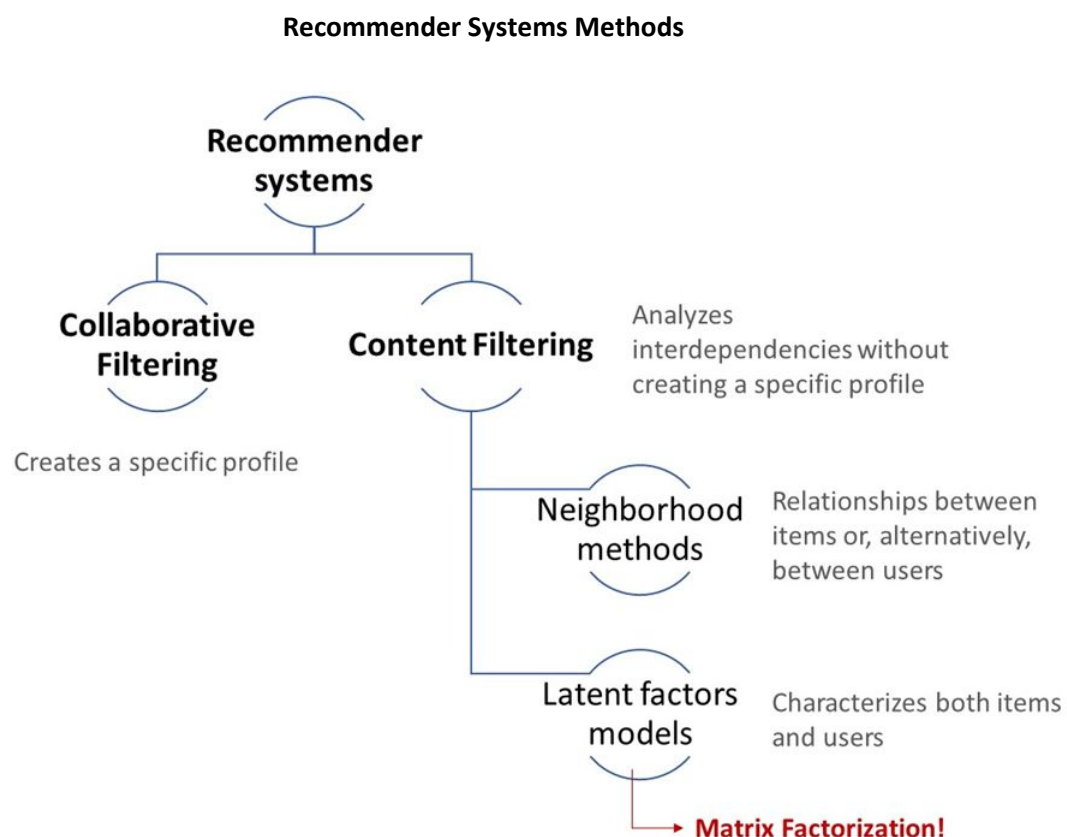
[https://datajobs.com/data-science-repo/Recommender-Systems-\[Netflix\].pdf](https://datajobs.com/data-science-repo/Recommender-Systems-[Netflix].pdf)

² Ensembling for the Netflix Prize http://web.stanford.edu/~lmackey/papers/netflix_story-nas11-slides.pdf

³ How the Netflix prize was won <http://blog.echen.me/2011/10/24/winning-the-netflix-prize-a-summary/>

Matrix Factorization belongs to one of two main strategies in recommender systems: *collaborative filtering*. In general, recommender systems are based either on *content filtering* (creates a profile for each user or product to characterize its nature), or collaborative filtering (relies only on past user behavior without requiring the creation of explicit profiles and analyzes relationships between users and interdependencies among products to identify new user-item associations).⁴ On its own, collaborative filtering has two main areas: *neighborhood methods* and *latent factor models*. Neighborhood methods focus on relationships between items or, alternatively, between users, while latent factor models focus on characterizing both items and users.

In other words, latent factor models transform both items and users to the same latent factor space. This latent space tries to explain ratings by characterizing both products and users on factors automatically inferred from user feedback.⁵ In this area, Matrix Factorization is one of the most popular methods.



A matrix factorization is performed with a singular value decomposition (SVD) on the ratings matrix (using stochastic gradient descent and regularizing the weights of the factors). Matrix factorization is particularly advantageous because it allows the incorporation of additional information. For instance, when explicit feedback is not available (ie. direct ratings on a business), one can infer user's


⁴ Matrix factorization for recommender systems,
[https://datajobs.com/data-science-repo/Recommender-Systems-\[Netflix\].pdf](https://datajobs.com/data-science-repo/Recommender-Systems-[Netflix].pdf)

⁵ Advances in Collaborative Filtering from the Netflix prize
<https://datajobs.com/data-science-repo/Collaborative-Filtering-%5BKoren-and-Bell%5D.pdf>

preferences using implicit feedback by observing behavior like purchase history, search patterns, etc.

⁶

Matrix Factorization model



Some SVD-inspired methods for matrix factorization include: **i) Standard SVD** (you can dot product the user's vectors with the movies or business' vectors to predict each individual rating); **ii) Asymmetric SVD** (instead of users having their own notion of factor vectors, we can represent users as a bag of items they have rated); **iii) SVD++** (incorporates both the standard SVD and the asymmetric SVD model by representing users both by their own factor representation and as a bag of item vectors)⁷. For purposes of this project, we will focus on the Standard SVD.

4. Modeling Approach and Project Trajectory

4.1. Baseline model

The baseline will be estimated with two models: Sample Averages and Regularized Regression. With the residuals of these models, we will estimate the matrix factorization that will allow to predict more accurately the recommendations.

a. Sample Averages

Suppose an user rates a restaurant with 4 stars⁸. This can be decomposed in:

- A baseline rating (average of all reviews) = 3.1.
- The effect of the user = - 0.5
- The effect of the restaurant = 0.8
- Specific interaction = 0.6

⁶ Matrix factorization for recommender systems,

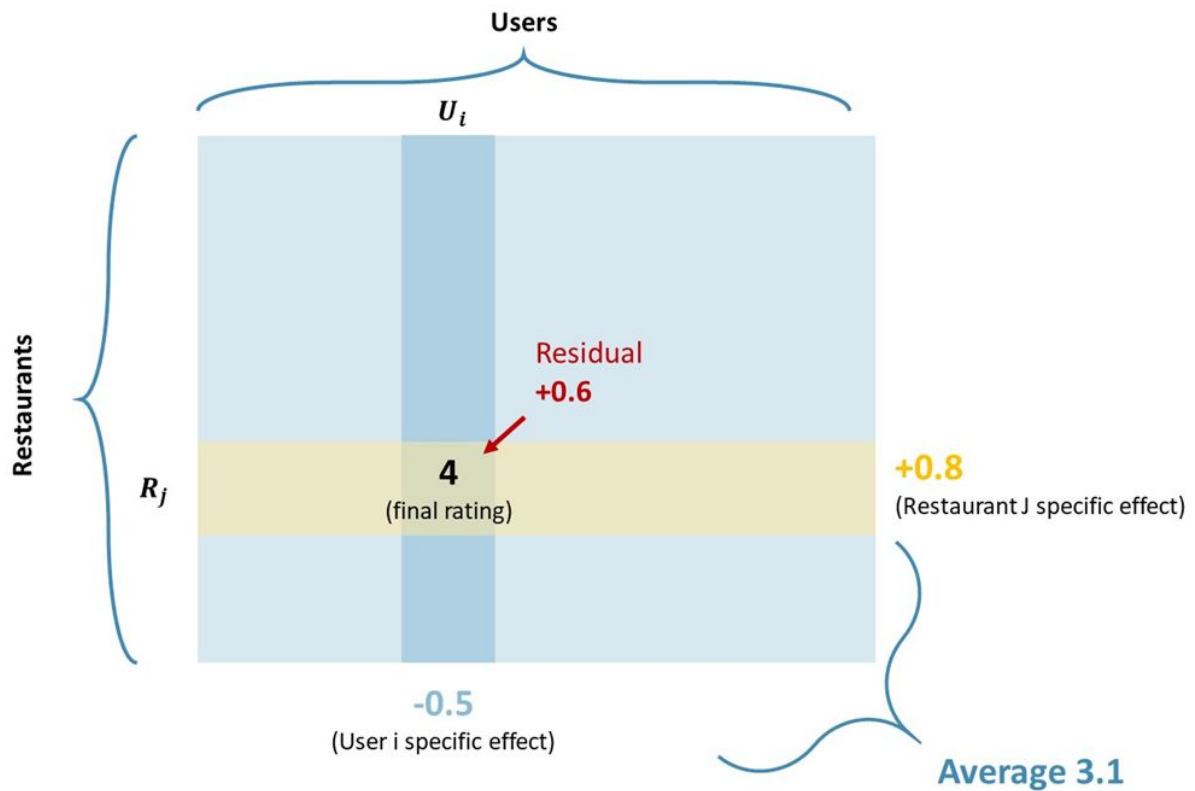
[https://datajobs.com/data-science-repo/Recommender-Systems-\[Netflix\].pdf](https://datajobs.com/data-science-repo/Recommender-Systems-[Netflix].pdf)

⁷ How the Netflix prize was won <http://blog.echen.me/2011/10/24/winning-the-netflix-prize-a-summary/>

⁸ Example inspired from How the Netflix prize was won

<http://blog.echen.me/2011/10/24/winning-the-netflix-prize-a-summary/>

Sample Averages Baseline Example



The effect of the specific interaction between the restaurant and the users is the residual, which will be estimated with the matrix factorization.

From the Yelp data, we estimated an average of 3.7 stars. With this information, we calculated the averages per user and restaurant, to predict the missing values.

The Root Mean Square Error of the predicted stars with the sample average baseline is **0.89** (in other words, less than one star of error).

b. Regularized Regression

Another way to estimate the baseline is through a regularized regression. To achieve this, we chose Lasso with cross validation (5 fold) to find the alpha that minimizes the Root Mean Square Error. The alpha was 0.001. This alpha gave us a RMSE of **0.935**.

c. Matrix Factorization

The matrix factorization was done with the residuals from the regularization regression. We predicted the residuals using the following parameters:

- $\lambda = 0.1$
- $n_{\text{factors}} = 20$
- $n_{\text{iterations}} = 10$

With the residuals, we estimated the final rating per review. Our final RMSE was **0.926**, less than the RMSE from the regression. However, this could have been optimized iterating over the parameters.

5. Results, Conclusions, and Future Work

As a result, the Factorization Matrix can be an effective way for Latent Factors Models, as it reduced the RMSE from the baseline regressions. Our results could further be improved by:

- Do a cross validation to choose the best parameters for the matrix factorization, including: number of iterations of ALS (`n_iterations`), learning rate (`lambda`) and number of latent factors (`n_factors`).
- Run the model on some virtual machine/on the cloud in order to account for the entire dataset, and not only users who gave 150+ reviews and restaurants with 350+ reviews.
- Do an ensemble model, taking into consideration the sample average baseline and the regression baseline. Also, adding biases and additional input sources can have a positive effect in the optimization.